

Google Tensor Processing Unit

- ▶ Google getting into hardware design. Three spins of TPU already!
- ▶ TPU is a specialized chip for AI \rightarrow matrix-matrix/tensor operations
- ▶ **Key Idea:** Systolic design
 - ▶ Lots of simple parallel computational units
 - ▶ Data movement strictly localized to immediate neighbours
- ▶ Tightly coupled with TensorFlow (programming APIs)
- ▶ \$6.50 USD per TPU per hour (as of today!)





Figure: TPU v3 device → liquid cooled!

<https://cloud.google.com/tpu/>

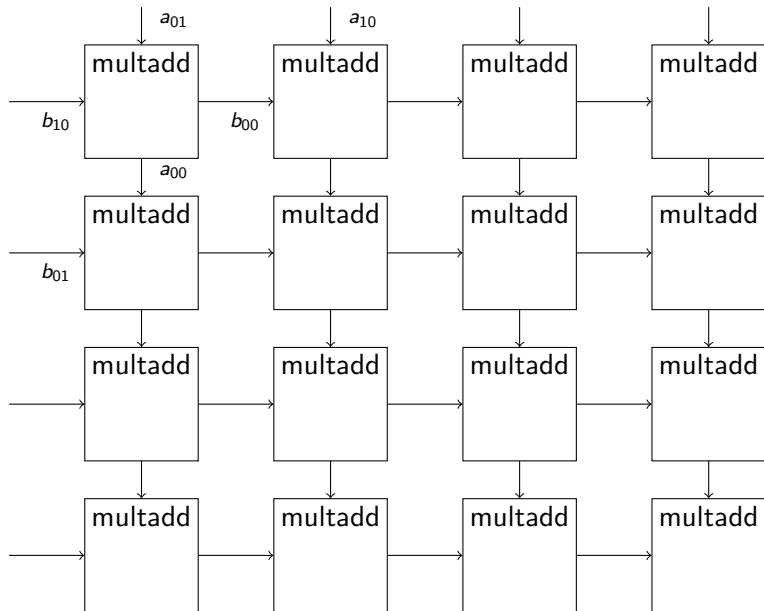
TPU History

- ▶ First generation
 - ▶ 64K 8-bit multiply-add operations @700 MHz \rightarrow 44 Tops/s
- ▶ Second generation
 - ▶ 32K 32-bit (and single-precision float) multiply-add operations @700 MHz \rightarrow 45 TFLOPS (float)
- ▶ Third generation
 - ▶ Not a lot of details \rightarrow $2\times$ better than TPU v2

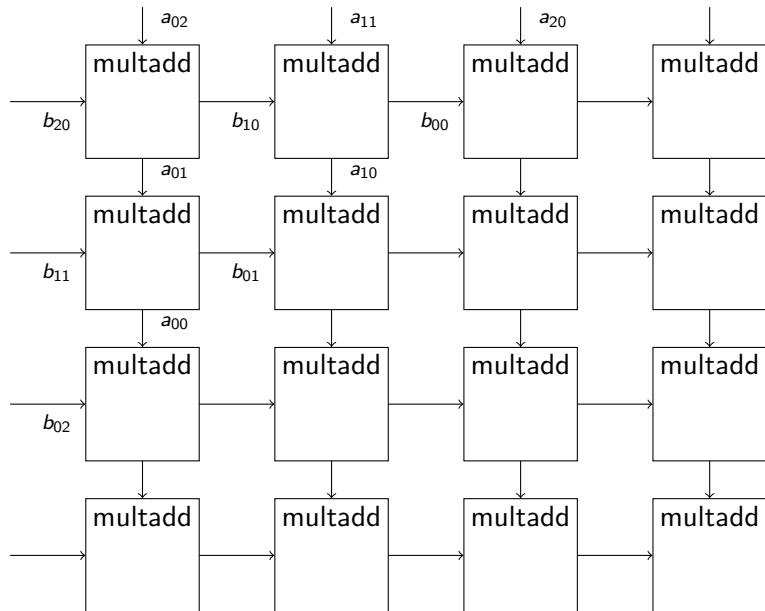
Systolic Array (Step-by-step Matrix Multiply)



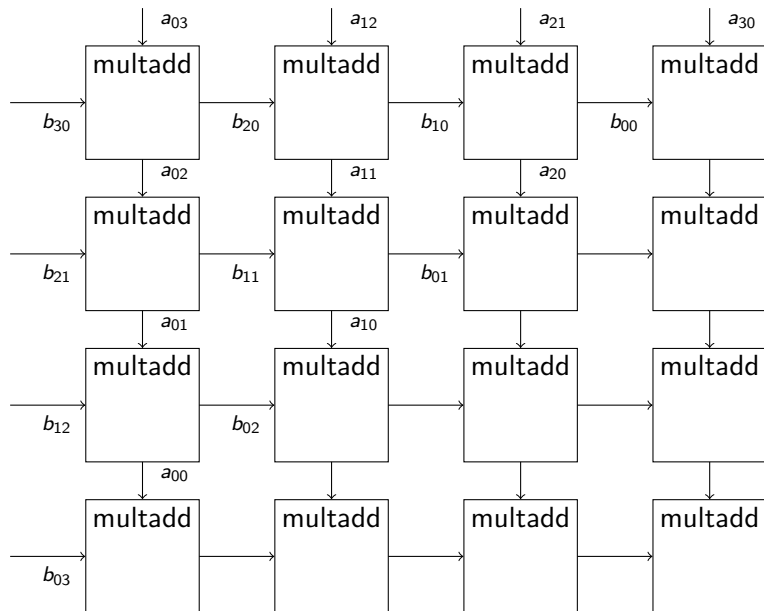
Systolic Array (Step-by-step Matrix Multiply)



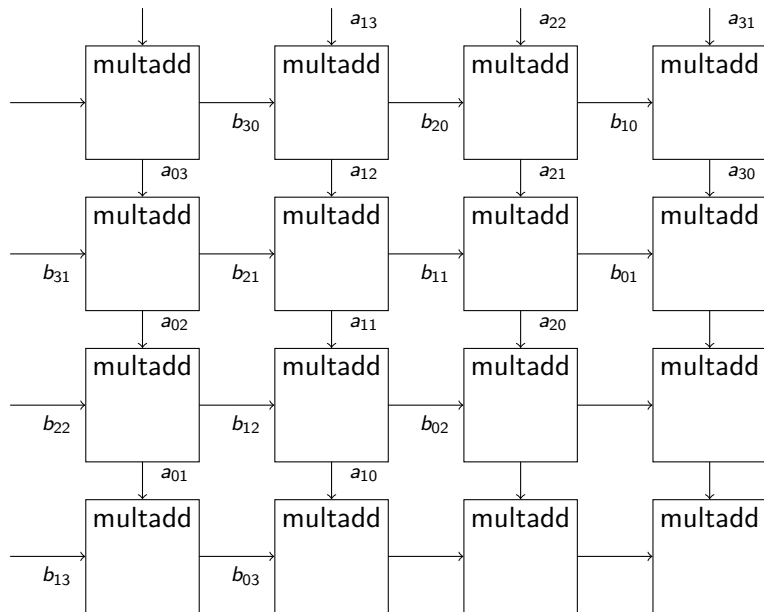
Systolic Array (Step-by-step Matrix Multiply)



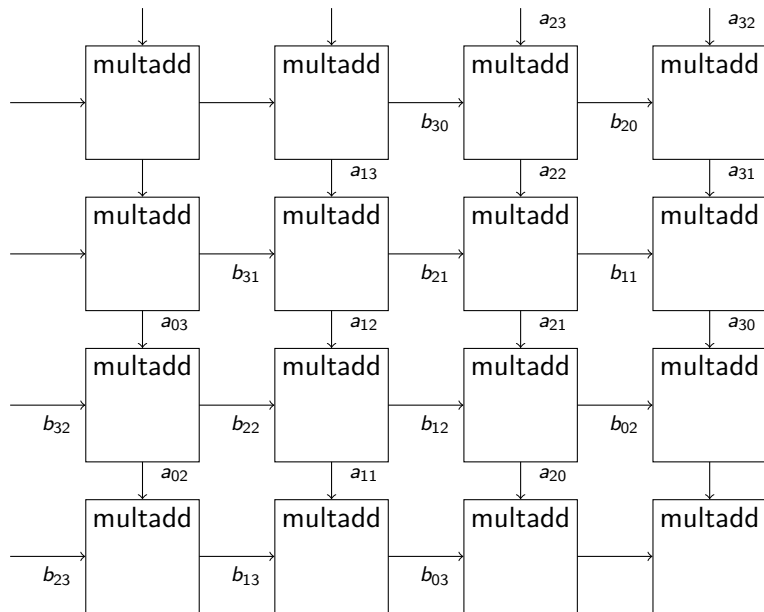
Systolic Array (Step-by-step Matrix Multiply)



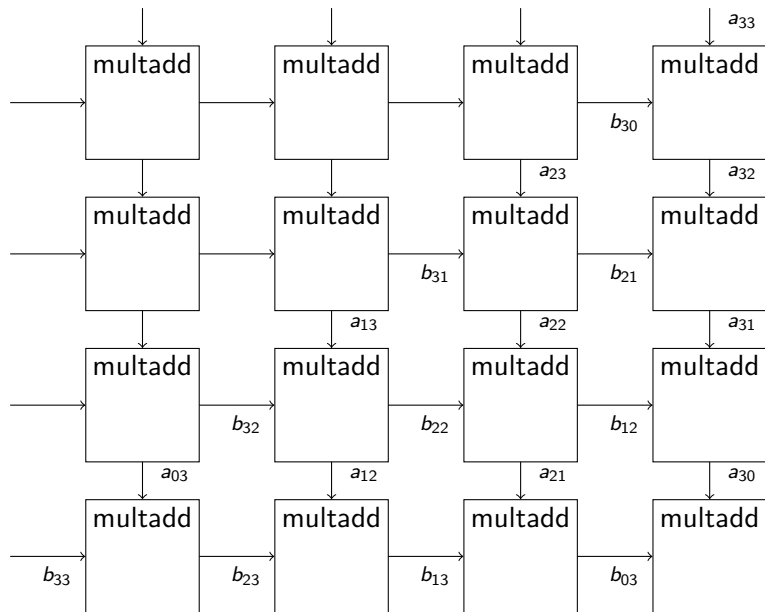
Systolic Array (Step-by-step Matrix Multiply)



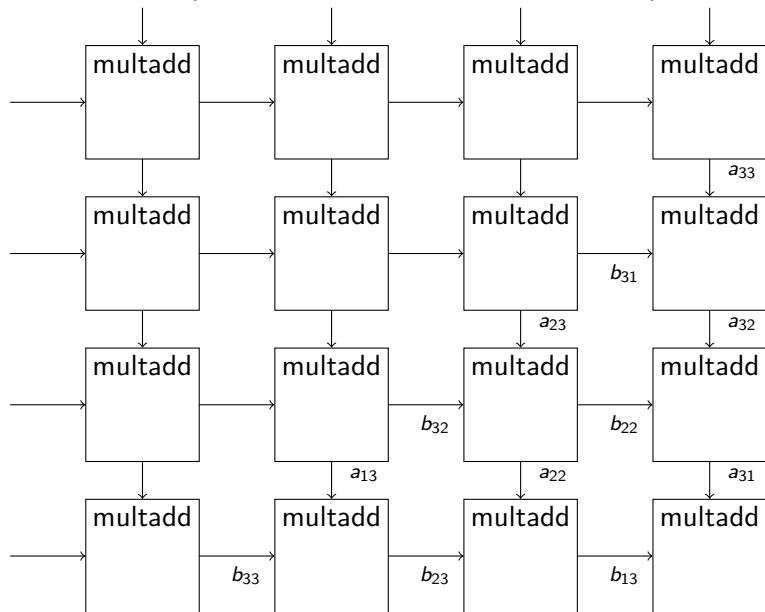
Systolic Array (Step-by-step Matrix Multiply)



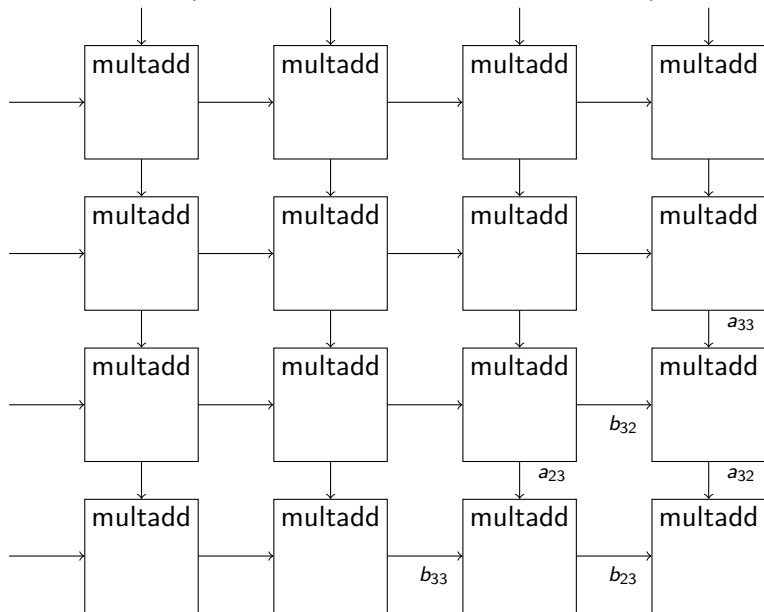
Systolic Array (Step-by-step Matrix Multiply)




Systolic Array (Step-by-step Matrix Multiply)




Systolic Array (Step-by-step Matrix Multiply)



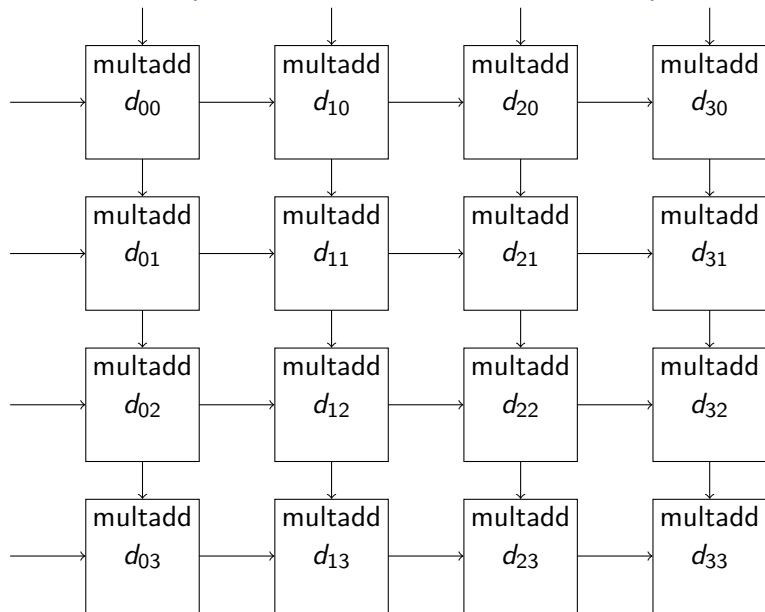
Systolic Array (Step-by-step Matrix Multiply)



Systolic Array (Step-by-step Matrix Multiply)





Systolic Array (Step-by-step Matrix Multiply)



Matrix-Matrix Multiplication on Systolic Array


- ▶ Send A matrix along column, and B matrix along rows of a 2D array
- ▶ Each step along the way, the compute block performs a small multiply-add operation
- ▶ Data movement is co-ordinated to ensure correct data arrives at correct compute block in the correct time!
- ▶ All data movement is local \rightarrow hardware design is super easy!
- ▶ In the end, final product D is available in every compute block \rightarrow result must be read out (not shown)

The multadd Building Block



- ▶ The multadd block is a simple $8b \times 8b$ multiplier followed by a 32b accumulator
- ▶ All communication in the chip is strictly local \rightarrow nearest neighbour \rightarrow no long wires!
 - ▶ Matrix elements *streamed* into systolic core from top+left links
 - ▶ Data forwarded to links below+right for further use by neighbouring elements
- ▶ The building block in the TPUv1 was just 8b multiplication of neuron weights, and input \rightarrow later revisions added floating-point and custom formats
- ▶ Design is so easy, even Google engineers can build this :)


RTL Description of the multadd block



```
module systolic_leaf
  (input clk,
   input rst,
   input [7:0] a_in,
   input [7:0] b_in,
   output wire [7:0] a_out,
   output wire [7:0] b_out
  );
```

TPUv1 has 8b multipliers, 32b accumulators


RTL Description of the multadd block



TPUv1 has 8b multipliers, 32b accumulators


```
reg [31:0] d;  
reg [7:0] a_tmp;  
reg [7:0] b_tmp;  
  
always@(posedge clk) begin  
    if (rst) begin  
        d    <= 0;  
        a_tmp <= 0;  
        b_tmp <= 0;  
    end else begin  
        d <= d + a_tmp*b_tmp;  
        a_tmp <= a_in;  
        b_tmp <= b_in;  
    end  
end  
  
assign a_out = a_tmp;  
assign b_out = b_tmp;  
  
endmodule
```

Systolic Array Design (4×4 design)




Caveat: Above pic is simple, actual TPUv1 is 256×256 blocks

RTL Description of the Systolic Array



```
module systolic_complete
# (parameter SIZE = 4)
(input clk,
 input rst,
 input wire [7:0] a [SIZE-1:0],
 input wire [7:0] b [SIZE-1:0]
);
```

RTL Description of the Systolic Array



```
genvar i,j;

wire [7:0] horizontal_wires [SIZE+1][SIZE];
wire [7:0] vertical_wires [SIZE][SIZE+1];

generate
for (i=0; i<=SIZE-1; i=i+1) begin : row
    assign vertical_wires[0][i] = b[i];
    assign horizontal_wires[i][0] = a[i];
    for (j=0; j<=SIZE-1; j=j+1) begin : col
        systolic_leaf
        systolic_leaf_inst (
            .clk      (clk),
            .rst      (rst),
            .a_in     (vertical_wires[i][j]),
            .b_in     (horizontal_wires[i][j]),
            .a_out    (vertical_wires[i+1][j]),
            .b_out    (horizontal_wires[i][j+1])
        );
    end
end
endgenerate

endmodule
```